

Derek Alyne (dalyne2)  
Joshua Sanchez (jsanch84)  
Howard Shan (howards2)  
ECE 408  
Professor Sanjay Patel  
8 March 2019

## Report

### I. Milestone 1

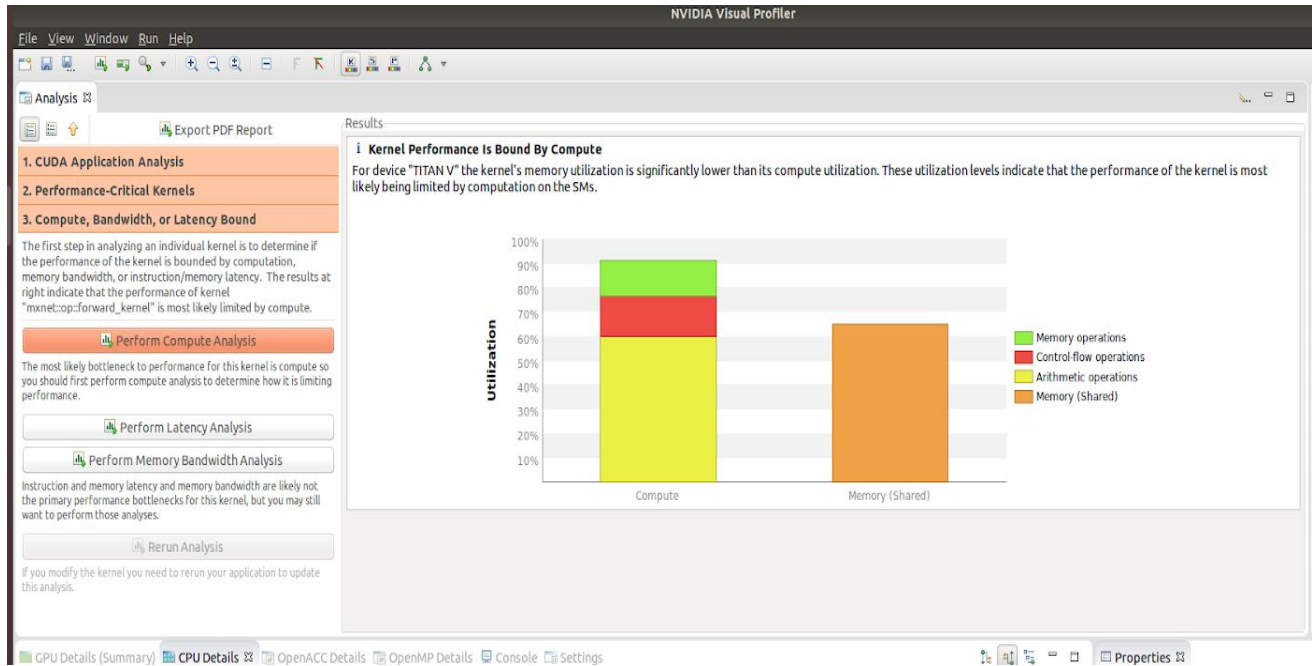
A. Include a list of all kernels that collectively consume more than 90% of the program time.

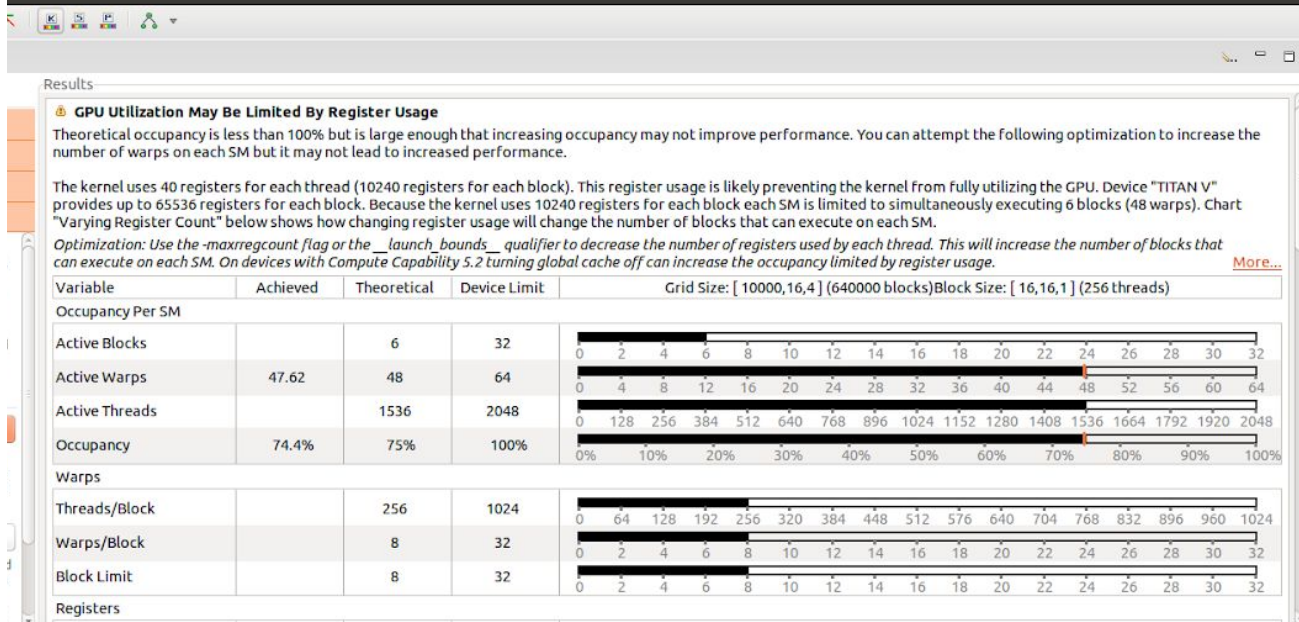
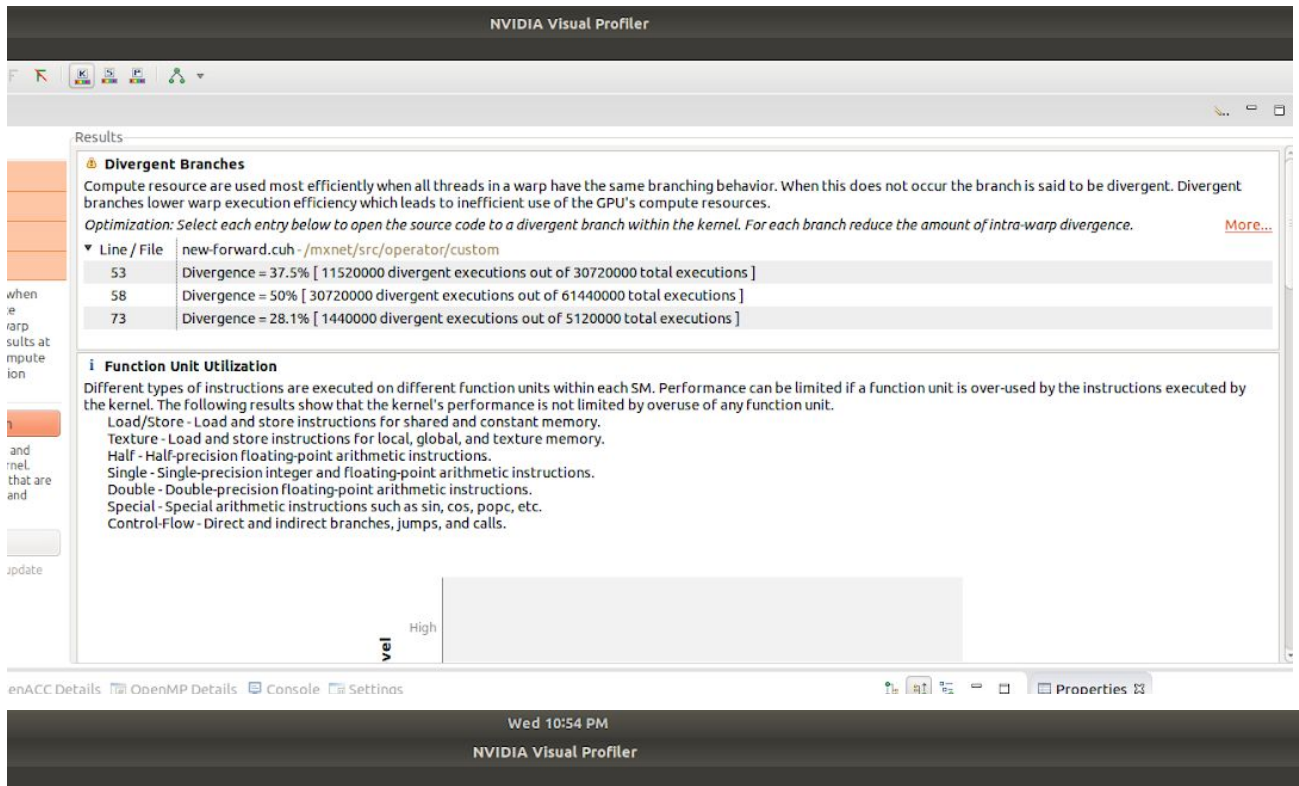
- [CUDA memcpy HtoD]
- `void cudnn::detail::implicit_convolve_sgemm<float, float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>(int, int, int, float const *, int, float*, cudnn::detail::implicit_convolve_sgemm<float, float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>*, kernel_conv_params, int, float, float, int, float, float, int, int)`
- `volta_cgemm_64x32_tn`
- `void op_generic_tensor_kernel<int=2, float, float, float, int=256, cudnnGenericOp_t=7, cudnnNanPropagation_t=0, cudnnDimOrder_t=0, int=1>(cudnnTensorStruct, float*, cudnnTensorStruct, float const *, cudnnTensorStruct, float const *, float, float, float, float, dimArray, reducedDivisorArray)`
- `void fft2d_c2r_32x32<float, bool=0, bool=0, unsigned int=1, bool=0, bool=0>(float*, float2 const *, int, int, int, int, int, int, int, float, float, cudnn::reduced_divisor, bool, float*, float*, int2, int, int)`
- `Volta_sgemm_128x128_tn`
- `void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0, bool=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)`


- `void fft2d_r2c_32x32<float, bool=0, unsigned int=0, bool=0>(float2*, float const *, int, int, int, int, int, int, int, int, int, cudnn::reduced_divisor, bool, int2, int, int)`
- B. Report: Include a list of all CUDA API calls that collectively consume more than 90% of the program time.
  - `cudaStreamCreateWithFlags`
  - `cudaMemGetInfo`
  - `cudaFree`
- C. Report: Include an explanation of the difference between kernels and API calls
  - API calls are functions defined by the CUDA library, such as `cudaMalloc` and `cudaMemcpy`, while kernels are functions that the programmer (or other libraries) defines to run on the gpu. They are a minimal set of extensions to the C language and a runtime library to help the programmer interface with the gpu. Kernel functions are typically run a large number of times in parallel, using multiple blocks and threads. According to the CUDA documentation, “A kernel is defined using the `__global__` declaration specifier and the number of CUDA threads that execute that kernel for a given kernel call is specified using a new `<<<...>>>`execution configuration syntax.”
- D. Report: Show output of rai running MXNet on the CPU
  - `* Running /usr/bin/time python ml.1.py`
  - `Loading fashion-mnist data... done`
  - `Loading model... done`
  - `New Inference`
  - `EvalMetric: {'accuracy': 0.8236}`
  - `8.91user 3.64system 0:05.11elapsed 245%CPU`  
`(0avgtext+0avgdata 2470716maxresident)k`
  - `0inp`
  - `uts+2824outputs (0major+666444minor)pagefaults 0swaps`
- E. Report: List program run time
  - 5.11 seconds
- F. Report: Show output of rai running MXNet on the GPU
  - `* Running /usr/bin/time python ml.2.py`
  - `Loading fashion-mnist data... done`
  - `Loading model... done`
  - `New Inference`
  - `EvalMetric: {'accuracy': 0.8236}`
  - `4.43user 3.37system 0:04.33elapsed 180%CPU`  
`(0avgtext+0avgdata 2841612maxresident)k`
  - `8inputs+1728outputs (0major+660933minor)pagefaults`  
`0swaps`
- G. Report: List program run time

- 4.33 seconds
- II. Milestone 2
  - A. Whole Program Execution Time
    - 12.13 seconds
  - B. OpTimes:
    - 2.583861 seconds
    - 7.785734 seconds
- III. Milestone 3

The following images are the result of loading analysis.nvvp into the Visual Profiler:







#### IV. Milestone 4